

Ivester Institute for Business Analytics and Insights: Intro to R for BUSN Analytics

Katherine A Ireland

2025-09-10

0) Background: Why use R for BUSN Analytics?

- Works well for the full data science pipeline: getting data, data wrangling, to analysis and visualization
- Consistent grammar → fast, repeatable visuals for reports/decks
- Clear mapping from business questions → visual encodings
- Works across disciplines + data

Quick references

- R for Data Science (R4DS), Chapter on Data Visualization (<https://r4ds.had.co.nz/data-visualisation.html>)
- dplyr documentation (<https://cran.r-project.org/web/packages/dplyr/vignettes/dplyr.html>)
- ggplot2 documentation (<https://ggplot2.tidyverse.org/>)
- Posit (RStudio) ggplot2 Cheatsheet (<https://posit.co/resources/cheatsheets/>)

the Gapminder dataset

- The Gapminder dataset (<https://www.gapminder.org/>) comes from the Gapminder Foundation, a non-profit started by Hans Rosling.
- It was created to make global development statistics accessible, visual, and easy to understand.
- The data combines information from sources like the United Nations, World Bank, and other international agencies.
- It's widely used in teaching because it's clean, tidy, and covers long-term trends in health and wealth across the world.

1) dplyr verbs: the powerhouse behind data processing and analytics

Next, we will go through several useful dplyr verbs one-by-one on our gapminder dataset.

Dplyr verbs operate on different aspects of your dataset: on rows (filter, slice, arrange), on columns (select, rename, mutate, relocate), on groups of rows (summarise)

filter() → Keep only certain rows.

What it does:

Think of `filter()` as a row *sieve*. You provide a condition, and only rows that meet the condition stay in the dataset.

Example: Get a specific row value.

```
# Example: filtering by two row values
gap %>%
  filter(continent == "Asia", year == 2007)
```

```
## # A tibble: 33 × 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      2007   43.8   31889923    975.
## 2 Bahrain     Asia      2007   75.6    708573    29796.
## 3 Bangladesh  Asia      2007   64.1  150448339   1391.
## 4 Cambodia    Asia      2007   59.7   14131858   1714.
## 5 China        Asia      2007   73.0  1318683096   4959.
## 6 Hong Kong, China Asia      2007   82.2    6980412   39725.
## 7 India        Asia      2007   64.7  1110396331   2452.
## 8 Indonesia    Asia      2007   70.6   223547000   3541.
## 9 Iran         Asia      2007   71.0   69453570   11606.
## 10 Iraq        Asia      2007   59.5   27499638    4471.
## # i 23 more rows
```

2. select() → Choose specific columns.

What it does:

Imagine a spreadsheet with many columns. `select()` lets you pick just the variables you care about (and reorder them if you want).

Example: Keep only five columns.

```
gap %>%
  select(country, continent, year, lifeExp, gdpPercap)
```

```
## # A tibble: 1,704 × 5
##   country      continent  year lifeExp gdpPercap
##   <fct>        <fct>    <int>   <dbl>   <dbl>
## 1 Afghanistan Asia      1952    28.8    779.
## 2 Afghanistan Asia      1957    30.3    821.
## 3 Afghanistan Asia      1962    32.0    853.
## 4 Afghanistan Asia      1967    34.0    836.
## 5 Afghanistan Asia      1972    36.1    740.
## 6 Afghanistan Asia      1977    38.4    786.
## 7 Afghanistan Asia      1982    39.9    978.
## 8 Afghanistan Asia      1987    40.8    852.
## 9 Afghanistan Asia      1992    41.7    649.
## 10 Afghanistan Asia      1997    41.8    635.
## # i 1,694 more rows
```

3. `arrange()` → Reorder rows.

What it does:

Think of `arrange()` as sorting your table. By default it sorts ascending (smallest → largest), but you can use `desc()` for descending order.

Example: sorting by population (ascending)

```
gap %>%
  arrange(pop) %>%    # sort ascending by population
  head(10)
```

```
## # A tibble: 10 × 6
##   country      continent  year lifeExp  pop gdpPercap
##   <fct>        <fct>    <int>   <dbl> <int>   <dbl>
## 1 Sao Tome and Principe Africa      1952    46.5 60011     880.
## 2 Sao Tome and Principe Africa      1957    48.9 61325     861.
## 3 Djibouti      Africa      1952    34.8 63149    2670.
## 4 Sao Tome and Principe Africa      1962    51.9 65345    1072.
## 5 Sao Tome and Principe Africa      1967    54.4 70787    1385.
## 6 Djibouti      Africa      1957    37.3 71851    2865.
## 7 Sao Tome and Principe Africa      1972    56.5 76595    1533.
## 8 Sao Tome and Principe Africa      1977    58.6 86796    1738.
## 9 Djibouti      Africa      1962    39.7 89898    3021.
## 10 Sao Tome and Principe Africa      1982    60.4 98593    1890.
```

```
# Next: get highest life expectancy (2007) using filter in combo w arrange
gap %>%
  filter(year == 2007) %>%
  arrange(desc(lifeExp)) %>%
  head(10)
```

```
## # A tibble: 10 × 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
## 1 Japan        Asia      2007    82.6 127467972  31656.
## 2 Hong Kong, China Asia      2007    82.2  6980412   39725.
## 3 Iceland      Europe    2007    81.8   301931   36181.
## 4 Switzerland Europe    2007    81.7  7554661   37506.
## 5 Australia    Oceania   2007    81.2 20434176  34435.
## 6 Spain        Europe    2007    80.9 40448191  28821.
## 7 Sweden       Europe    2007    80.9  9031088   33860.
## 8 Israel       Asia      2007    80.7  6426679   25523.
## 9 France       Europe    2007    80.7 61083916  30470.
## 10 Canada      Americas  2007    80.7 33390141  36319.
```

4. mutate create new columns.

What it does: adds new variables (columns) or transforms existing ones.

- You can base the new variable on any calculation, combination, or condition from existing columns.

Example: compare gdp across countries

```

#first get gdp from gdp per capita
gap2<- gap %>%mutate(gdp = gdpPercap * pop)    # total GDP
#View(gap2)

#scale in millions
gap2<- gap %>%
  mutate(gdp_million = (gdpPercap * pop) / 1e6)
#add in additional filtering
Arg_1982 <- gap2 %>% filter(country == "Argentina", year == 1982) %>%
  select(country, year, gdp_million)

##your turn! Let's try adding in this filter with mutate (filter(year == 2007, countr
y %in% c("China", "India", "United States", "Luxembourg"))) to compare gdp results ac
ross these countries

##another mutate example: what's going on here?
gap %>%
  mutate(pop_millions = pop / 1e6) %>%
  select(country, year, pop, pop_millions) %>%
  head(5)

```

```

## # A tibble: 5 × 4
##   country      year      pop pop_millions
##   <fct>      <int>    <int>      <dbl>
## 1 Afghanistan 1952  8425333      8.43
## 2 Afghanistan 1957  9240934      9.24
## 3 Afghanistan 1962 10267083     10.3
## 4 Afghanistan 1967 11537966     11.5
## 5 Afghanistan 1972 13079460     13.1

```

```

##and here?
gap %>%
  group_by(country) %>%
  mutate(is_growing = pop > lag(pop)) %>%
  select(country, year, pop, is_growing) %>%
  filter(country == "China") %>%
  head(10)

```

```
## # A tibble: 10 × 4
## # Groups:   country [1]
##   country year      pop is_growing
##   <fct>   <int>    <int> <lgl>
## 1 China   1952  556263527 NA
## 2 China   1957  637408000 TRUE
## 3 China   1962  665770000 TRUE
## 4 China   1967  754550000 TRUE
## 5 China   1972  862030000 TRUE
## 6 China   1977  943455000 TRUE
## 7 China   1982 1000281000 TRUE
## 8 China   1987 1084035000 TRUE
## 9 China   1992 1164970000 TRUE
## 10 China  1997 1230075000 TRUE
```

group_by and summarise.

What it does: Summarises data into groups by calculating values like mean, sum, count, etc.

Example: compare gdp across countries.

- First use `group_by()` to define categories.
- use `summarise()` to collapse each group into one row with summary statistics.

```
# Average life expectancy by continent in 2007
gap %>%
  filter(year == 2007) %>%
  group_by(continent) %>%
  summarise(avg_lifeExp = mean(lifeExp), .groups = "drop") %>%
  arrange(desc(avg_lifeExp))
```

```
## # A tibble: 5 × 2
##   continent avg_lifeExp
##   <fct>      <dbl>
## 1 Oceania      80.7
## 2 Europe       77.6
## 3 Americas     73.6
## 4 Asia         70.7
## 5 Africa       54.8
```

count

What it does: generates quick frequencies of data - Quickly counts the number of rows in each category.

Example: number of rows per continent (sorted).

```
gap %>% count(continent, sort = TRUE)
```

```
## # A tibble: 5 × 2
##   continent      n
##   <fct>        <int>
## 1 Africa        624
## 2 Asia          396
## 3 Europe        360
## 4 Americas      300
## 5 Oceania       24
```

slice_max→ Get the “top N” rows

What it does: Selects the highest (or lowest - with slice_min()) values within a column.

Example: top 5 countries by GDP per capita in 2007

```
# Top 5 countries by GDP per capita in 2007
gap %>%
  filter(year == 2007) %>%
  slice_max(gdpPercap, n = 5, with_ties = FALSE) %>%
  select(country, continent, gdpPercap)
```

```
## # A tibble: 5 × 3
##   country      continent gdpPercap
##   <fct>        <fct>        <dbl>
## 1 Norway      Europe        49357.
## 2 Kuwait      Asia          47307.
## 3 Singapore   Asia          47143.
## 4 United States Americas      42952.
## 5 Ireland     Europe        40676.
```


now in groups of two or three answer this question using our dplyr verbs

what top 2 countries per continent had the highest gdp per capita in 1982 and in 2007?


```
top2 <- gap %>% filter(year == 1982) %>%
  group_by(continent) %>%
  slice_max(order_by = gdpPercap, n = 2) %>%
  select(country, gdpPercap)
#View(top2)
```

Visualization w ggplot

```
# Bar plot
top2 %>%
  ggplot(aes(x = reorder(country, gdpPercap), y = gdpPercap, fill = continent)) +
  geom_col(show.legend = FALSE) +
  coord_flip() +
  labs(
    title = "Top 2 Countries per Continent by GDP per Capita (1982)",
    x = NULL,
    y = "GDP per Capita (USD)"
  )
)
```




```
#compare distribution of life expectancies across continents in 1982
gap %>%
  filter(year == 1982) %>%
  ggplot(aes(continent, lifeExp)) +
  geom_boxplot()
```

```
#correlation: visualize the relationship btw wealth and life expectancy
gap %>%
  filter(year == 2007) %>%
  summarise(cor(gdpPercap, lifeExp))
```

```
## # A tibble: 1 × 1
##   `cor(gdpPercap, lifeExp)`
##   <dbl>
## 1 0.679
```

```
gap %>%
  filter(year == 2007) %>%
  ggplot(aes(x = gdpPercap, y = lifeExp, color = continent, size = pop)) +
  geom_point(alpha = 0.7) +
  scale_x_log10() + #applies log base so easier to see lower income
  labs(
    title = "GDP per Capita vs Life Expectancy (1982)",
    x = "GDP per Capita (log scale, USD)",
    y = "Life Expectancy (years)"
  ) +
  guides(size = "none")
```



```
#trends over time
gap %>%
  group_by(continent, year) %>%
  summarise(avg_lifeExp = mean(lifeExp), .groups = "drop") %>%
  ggplot(aes(year, avg_lifeExp, color = continent)) +
  geom_line()
```

